

# Machine Learning 2023/2024 (2<sup>nd</sup> semester)



Master in Electrical and Computer Engineering

Department of Electrical and Computer Engineering

**A. Pedro Aguiar** ([pedro.aguiar@fe.up.pt](mailto:pedro.aguiar@fe.up.pt)), **Aníbal Matos** ([anibal@fe.up.pt](mailto:anibal@fe.up.pt)), **Andry Pinto** ([amgp@fe.up.pt](mailto:amgp@fe.up.pt)), **Daniel Campos** ([dfcampos@fe.up.pt](mailto:dfcampos@fe.up.pt)), **Maria Inês Pereira** ([maria.ines@fe.up.pt](mailto:maria.ines@fe.up.pt))

FEUP, Mar. 2024

---

## Notebook 04

- ✓ Part 1 - Linear Regression with regularization
- ✓ The Generalization problem

Understanding the following challenges is crucial for machine learning practitioners:

1. How does the model perform on ('training') data which has experienced before?
2. How does a model perform on previously unseen ('testing') data?
3. How should a model be selected?

Those questions constitute the basics of **underfitting** and **overfitting**.

### Example:

Let's produce noisy data that will be fitted later on by different models (with polynomial and Gaussian kernels).

```
%matplotlib inline
import matplotlib.pyplot as plt
import seaborn as sns; sns.set()
import numpy as np
from scipy import linalg
from sklearn.metrics import mean_squared_error
from sklearn.linear_model import LinearRegression
```

```

from sklearn.pipeline import make_pipeline
from sklearn.preprocessing import PolynomialFeatures

rng = np.random.RandomState(1)

size_training = 50
x_limit = 8

#Create some noisy data that will be used for training models
x_training = x_limit * rng.rand(size_training)
y_training = np.sin(x_training+3.1415/2) + 0.1 * rng.randn(size_training) # s

plt.scatter(x_training, y_training, marker="x")
plt.title("Training set");

```



1) Using models based on **Polynomial Kernels** to fit this training set.

```

# Using sklearn
from sklearn import linear_model
from sklearn.preprocessing import PolynomialFeatures
from sklearn.metrics import mean_squared_error, r2_score

```

```
#testing set
size_testing = 20
x_testing = np.linspace(0, x_limit, size_testing)
y_testing = np.sin(x_testing+3.1415/2) + 0.1 * rng.randn(size_testing)

# continous model (for drawing)
draw_size = 500
x_draw = np.linspace(0,x_limit,draw_size)

#Overfitting Model
nb_degree_over = 12
model_overfitting = make_pipeline(PolynomialFeatures(nb_degree_over), Linear
model_overfitting.fit(x_training[:, np.newaxis], y_training)
print("Training MSE (Poly Model Overfitting):", mean_squared_error(y_training,

y_overfitting = model_overfitting.predict(x_testing[:, np.newaxis])
m_error_overfitting = mean_squared_error(y_testing, y_overfitting)
print("Testing MSE (Poly Model Overfitting):", m_error_overfitting)

y_over_draw = model_overfitting.predict(x_draw[:, np.newaxis])

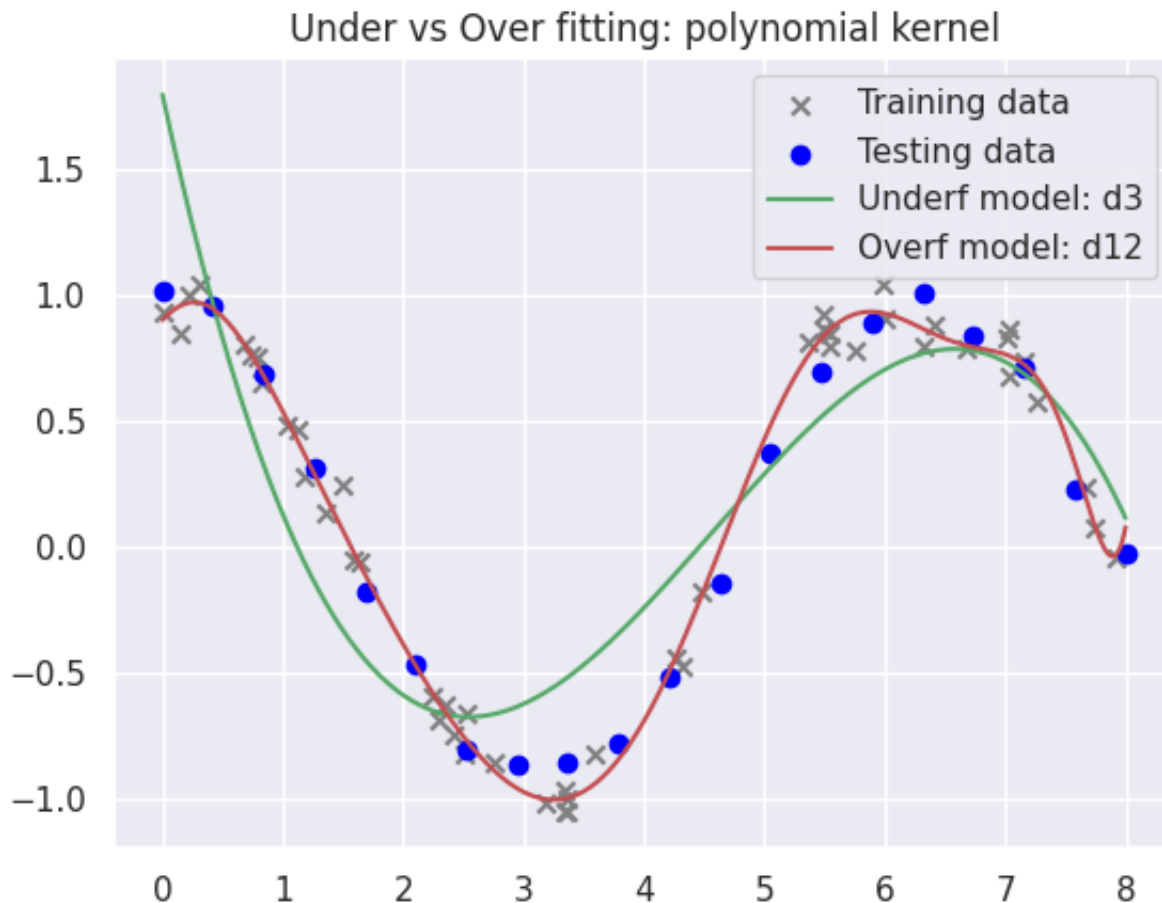
#Underfitting Model
nb_degree_under = 3
model_underfitting = make_pipeline(PolynomialFeatures(nb_degree_under), Line
model_underfitting.fit(x_training[:, np.newaxis], y_training)
print("Training MSE (Poly Model Underfitting):", mean_squared_error(y_training,

y_underfitting = model_underfitting.predict(x_testing[:, np.newaxis])
m_error_underfitting = mean_squared_error(y_testing, y_underfitting)
print("Testing MSE (Poly Model Underfitting):", m_error_underfitting)

y_under_draw = model_underfitting.predict(x_draw[:, np.newaxis])

#Plots
legend_fitting_model_overfitting = 'Overf model: d{}'.format(nb_degree_over)
legend_fitting_model_underfitting = 'Underf model: d{}'.format(nb_degree_under)
plt.scatter(x_training, y_training, color="gray", marker="x", label='Training
plt.scatter(x_testing, y_testing, color="blue", marker="o", label='Testing dat
plt.plot(x_draw, y_under_draw, color="g", label=legend_fitting_model_underfitti
plt.plot(x_draw, y_over_draw, color="r", label=legend_fitting_model_overfitting
plt.legend(loc ="best")
plt.title("Under vs Over fitting: polynomial kernel")
plt.show()
```

Training MSE (Poly Model Overfitting): 0.0042804398954628745  
 Testing MSE (Poly Model Overfitting): 0.007293292125722  
 Training MSE (Poly Model Underfitting): 0.11017920596696776  
 Testing MSE (Poly Model Underfitting): 0.09202245464570646



2) Using models based on **Gaussian Kernels** to fit this training set.

```

from sklearn.base import BaseEstimator, TransformerMixin
"""
    Gaussian models
    -> Uniformly spaced
    -> One-dimension input
"""
class GaussianBases(BaseEstimator, TransformerMixin):
    def __init__(self, N, width_kernel = 3.0):
        self.N = N
        self.width_kernel = width_kernel

    @staticmethod
    def _gauss_basis(x, y, width = 3.0, sigma_f=1, axis=None):
        arg = (x - y) / width
        return sigma_f*np.exp(-0.5 * np.sum(arg ** 2, axis))/(2 * width**2) ##
  
```

```

def fit(self, X, y=None):
    self.centers_ = np.linspace(X.min(), X.max(), self.N)
    self.width_ = self.width_kernel * (self.centers_[1] - self.centers_[0])
    return self

def transform(self, X):
    return self._gauss_basis(X[:, :, np.newaxis], self.centers_, self.width_)

#Overfitting Model
n_models_gauss_over = 20
gauss_model_overfitting = make_pipeline(GaussianBases(n_models_gauss_over), LinearRegression())
gauss_model_overfitting.fit(x_training[:, np.newaxis], y_training)
print("Training MSE (Gauss Model Overfitting):", mean_squared_error(y_training,
yfit_overfitting = gauss_model_overfitting.predict(x_testing[:, np.newaxis]
m_error_overfitting = mean_squared_error(y_testing, yfit_overfitting)
print("Testing MSE (Gauss Model Overfitting):", m_error_overfitting)

ydraw_over = gauss_model_overfitting.predict(x_draw[:, np.newaxis])


#Underfitting Model
n_models_gauss_under = 2
gauss_model_underfitting = make_pipeline(GaussianBases(3), LinearRegression())
gauss_model_underfitting.fit(x_training[:, np.newaxis], y_training)
print("Training MSE (Gauss Model Underfitting):", mean_squared_error(y_training,
yfit_underfitting = gauss_model_underfitting.predict(x_testing[:, np.newaxis]
m_error_underfitting = mean_squared_error(y_testing, yfit_underfitting)
print("Testing MSE (Gauss Model Underfitting):", m_error_underfitting)

ydraw_under = gauss_model_underfitting.predict(x_draw[:, np.newaxis])

#Plots
legend_fitting_model_overfitting = 'Overf model: n{}'.format(n_models_gauss_over)
legend_fitting_model_underfitting = 'Underf model: n{}'.format(n_models_gauss_under)
plt.scatter(x_training, y_training, color="gray", marker="x", label='Training data')
plt.scatter(x_testing, y_testing, color="blue", marker="o", label='Testing data')
plt.plot(x_draw, ydraw_under, color="g", label=legend_fitting_model_underfitting)
plt.plot(x_draw, ydraw_over, color="r", label=legend_fitting_model_overfitting)
plt.legend(loc="best")
plt.title("Under vs Over fitting: Gaussian kernel")
plt.show()


```

Training MSE (Gauss Model Overfitting): 0.0027875487865052385  
Testing MSE (Gauss Model Overfitting): 0.04623096259514714  
Training MSE (Gauss Model Underfitting): 0.23090736685120572  
Testing MSE (Gauss Model Underfitting): 0.21898176325660207



**Underfitting** means that the model is too simple and there is still progress to be made because the training error (considering a loss function) is high well after training the model with several examples.

**Overfitting** means that the model is fitting the training data more closely than the underlying phenomenon that is being captured and, therefore, there is a large difference between training (small) and testing error (high). It usually appears when the model fits the training data much better than new test examples - lack of generalization capacity.



source: (Goodfellow et al., 2016)

A model that is highly flexible will learn spurious (or irrelevant) patterns as easily as the true associations. Thus, the model will be unable to generalize well for unseen data. The best option to prevent a model from learning misleading patterns is to get more training data (which can be problematic for some cases). There are techniques that can be used to combat overfitting - a high model complexity.

## ✓ Regularization

A **Regularization** technique aims to reduce overfitting by penalizing the **model complexity**. Recalling that the linear regression minimizes the following loss function:

$$J(\theta) = \sum_{n=1}^N (y_n - \theta^T \mathbf{x}_n)^2$$

The searching space of  $\theta$  can be restricted by controlling  $\|\theta\|^2$  or  $\|\theta\|_1$ .



## ✓ Ridge regression ( $\ell_2$ Regularization)

The *Ridge regression* is a  $\ell_2$  regularization that penalizes the sum of squares ( $\ell_2$ -norms) of the model coefficients:

$$J_\lambda(\theta) = \sum_{n=1}^N (y_n - \theta^T \mathbf{x}_n)^2 + \lambda \|\theta\|^2$$

where  $\lambda \geq 0$  is a free parameter that quantifies a penalty for the norm of the parameter vector. Ridge regression tries to keep the sum of squared errors small and, at the same time, it attempts to reduce the norm of the estimated vector. This type of penalized model is built into sklearn with the `Ridge` estimator:

```
from sklearn.linear_model import Ridge


#Ridge Model
gauss_model_ridge = make_pipeline(GaussianBases(n_models_gauss_over), Ridge(alpha=1)
gauss_model_ridge.fit(x_training[:, np.newaxis], y_training)
print("Training MSE (Ridge Model):", mean_squared_error(y_training, gauss_model_ridge.predict(x_training[:, np.newaxis])))
print("Testing MSE (Ridge Model):", mean_squared_error(y_testing, gauss_model_ridge.predict(x_testing[:, np.newaxis])))

ydraw_ridge = gauss_model_ridge.predict(x_draw[:, np.newaxis])

#Plots
legend_fitting_model_ridge = 'Ridge model: n{}'.format(n_models_gauss_over)
plt.scatter(x_training, y_training, color="gray", marker="x", label='Training data')
plt.scatter(x_testing, y_testing, color="blue", marker="o", label='Testing data')
plt.plot(x_draw, ydraw_under, color="g", label=legend_fitting_model_underfitting)
plt.plot(x_draw, ydraw_over, color="r", label=legend_fitting_model_overfitting)
plt.plot(x_draw, ydraw_ridge, color="y", label=legend_fitting_model_ridge, linecolor='black')
plt.legend(loc="best")
plt.title("Ridge Regression: Gaussian kernel")
plt.show()
```

Training MSE (Ridge Model): 0.010903318981812961

Testing MSE (Ridge Model): 0.010359667098565731



## ✓ LASSO regression ( $\ell_1$ regularization)

The LASSO regularization involves penalizing the sum of absolute values ( $\ell_1$ -norm) of regression coefficients:

$$J_{\lambda}(\theta) = \sum_{n=1}^N (y_n - \theta^T \mathbf{x}_n)^2 + \lambda \|\theta\|_1$$

where  $\|\cdot\|_1$  is the  $\ell_1$ -norm (that is,  $\|\theta\|_1 = \sum_{i=1}^d \theta_i$ ) and  $\lambda \geq 0$  is a hyperparameter that tunes the strength of the penalty, and should be determined via, for example, cross-validation. Lasso regression tends to favor *sparse models* where possible: that is, it preferentially sets model coefficients to exactly zero. This type of penalized model is built into sklearn with the Lasso estimator:

```
from sklearn.linear_model import Lasso
```

```
#Lasso Model
```

```

gauss_model_lasso = make_pipeline(GaussianBases(n_models_gauss_over), Lasso(alpha=0.01))
gauss_model_lasso.fit(x_training[:, np.newaxis], y_training)
print("Training MSE (LASSO Model):", mean_squared_error(y_training, gauss_model_lasso.predict(x_training[:, np.newaxis])))
print("Testing MSE (LASSO Model):", mean_squared_error(y_testing, gauss_model_lasso.predict(x_testing[:, np.newaxis])))

ydraw_lasso = gauss_model_lasso.predict(x_draw[:, np.newaxis])

```

#Plots


```

legend_fitting_model_lasso = 'Lasso model: n{}'.format(n_models_gauss_over)
plt.scatter(x_training, y_training, color="gray", marker="x", label='Training data')
plt.scatter(x_testing, y_testing, color="blue", marker="o", label='Testing data')
plt.plot(x_draw, ydraw_under, color="g", label=legend_fitting_model_underfitting)
plt.plot(x_draw, ydraw_over, color="r", label=legend_fitting_model_overfitting)
plt.plot(x_draw, ydraw_ridge, color="y", label=legend_fitting_model_ridge)
plt.plot(x_draw, ydraw_lasso, color="m", label=legend_fitting_model_lasso, line)
plt.legend(loc="best")
plt.title("LASSO Regression: Gaussian kernel")
plt.show()

```

Training MSE (LASSO Model): 0.010903318981812961

Testing MSE (LASSO Model): 0.0074394158322638555



The estimated vector  $\theta$  for the Linear Regression, Ridge and Lasso models can be seen below. The majority of the coefficients are exactly zero with the LASSO regression penalty.

```
print("Model (Overfit):\n", gauss_model_overfitting.named_steps['linearregressi
print("Ridge Model:\n", gauss_model_ridge.named_steps['ridge'].coef_)
print("LASSO Model:\n", gauss_model_lasso.named_steps['lasso'].coef_)
```

Model (Overfit):

```
[-3.717e+07  2.627e+08 -9.827e+08  2.549e+09 -5.094e+09  8.261e+09
-1.116e+10  1.267e+10 -1.195e+10  8.950e+09 -4.489e+09 -6.236e+07
 3.402e+09 -4.853e+09  4.564e+09 -3.270e+09  1.818e+09 -7.584e+08
 2.157e+08 -3.212e+07]
```

Ridge Model:

```
[ 1.004  0.917  0.674  0.317 -0.089 -0.477 -0.794 -0.996 -1.041 -0.899
-0.572 -0.115  0.368  0.753  0.937  0.873  0.582  0.149 -0.306 -0.664]
```

LASSO Model:

```
[ 3.035  0.531  0.      0.      0.     -0.     -0.     -0.     -3.173 -0.
-0.     -0.     0.      0.      3.156  0.699  0.      0.     -0.     -0.   ]
```

## ✓ Activity 1.1

In the example above test several penalties for the Ridge and LASSO regressions and compare the mean square errors of training and testing sets.

For Ridge test with  $\lambda=0.01$ ,  $\lambda=0.1$ ,  $\lambda=1$  and  $\lambda=10$ .

For Lasso test with  $\lambda=0.0001$ ,  $\lambda=0.001$ ,  $\lambda=0.01$  and  $\lambda=0.1$ .

```
# code for activity 1.1
```

```
# solution 1.1
```

```
#Ridge Model
```

```
for a in [0.01, 0.1, 1, 10] :
    print("(Ridge) alpha=", a)
    gauss_model_ridge = make_pipeline(GaussianBases(n_models_gauss_over), Ridge(a))
    gauss_model_ridge.fit(x_training[:, np.newaxis], y_training)
    print("  Training MSE:", mean_squared_error(y_training, gauss_model_ridge.predict(x_training)))
    print("  Testing MSE :", mean_squared_error(y_testing, gauss_model_ridge.predict(x_testing)))
```

```
#Lasso Model
```

```
for a in [0.0001, 0.001, 0.01, 0.1] :
    print("(LASSO) alpha=", a)
    gauss_model_lasso = make_pipeline(GaussianBases(n_models_gauss_over), Lasso(a))
    gauss_model_lasso.fit(x_training[:, np.newaxis], y_training)
    print("  Training MSE:", mean_squared_error(y_training, gauss_model_lasso.predict(x_training)))
    print("  Testing MSE :", mean_squared_error(y_testing, gauss_model_lasso.predict(x_testing)))
```

```
(Ridge) alpha= 0.01
  Training MSE: 0.008053237966544727
  Testing MSE : 0.005462014355738859
(Ridge) alpha= 0.1
  Training MSE: 0.010903318981812961
  Testing MSE : 0.010359667098565731
(Ridge) alpha= 1
  Training MSE: 0.05632997997047238
  Testing MSE : 0.05073004195505988
(Ridge) alpha= 10
  Training MSE: 0.29351449910627253
  Testing MSE : 0.25773242698292087
(LASSO) alpha= 0.0001
  Training MSE: 0.007488561678321872
  Testing MSE : 0.004455933922947243
(LASSO) alpha= 0.001
  Training MSE: 0.008914237356444752
  Testing MSE : 0.0074394158322638555
(LASSO) alpha= 0.01
  Training MSE: 0.06568650481471255
  Testing MSE : 0.07355361298863375
(LASSO) alpha= 0.1
  Training MSE: 0.5114696052079297
  Testing MSE : 0.4522947801076927
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_coordinate_descent.py:111:
model = cd_fast.enet_coordinate_descent(
```

## ✓ Activity 1.2

Use the classic [Auto MPG dataset](#) to build models for predicting the fuel efficiency of cars from the late-1970s and early 1980s. Before starting, it is crucial to preprocess the data because contains missing information, presence of categorical values, and features at different scales. Please analyse and execute the following code that preprocesses data accordingly.

```
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import seaborn as sns
# Make NumPy printouts easier to read.
np.set_printoptions(precision=3, suppress=True)

# Download dataset
url = 'http://archive.ics.uci.edu/ml/machine-learning-databases/auto-mpg/auto-mpg.data'
column_names = ['MPG', 'Cylinders', 'Displacement', 'Horsepower', 'Weight',
                'Acceleration', 'Model Year', 'Origin']

raw_dataset = pd.read_csv(url, names=column_names,
                          na_values='?', comment='\t',
                          sep=' ', skipinitialspace=True)

dataset = raw_dataset.copy()

# Drop NaN values
dataset = dataset.dropna()

# One-hot encoding for "Origin" column, which is categorical and not numeric.
dataset['Origin'] = dataset['Origin'].map({1: 'USA', 2: 'Europe', 3: 'Japan'})
dataset = pd.get_dummies(dataset, columns=['Origin'], prefix='', prefix_sep='')

dataset
#dataset.tail()
```

	MPG	Cylinders	Displacement	Horsepower	Weight	Acceleration	Model Year	Europe
0	18.0	8	307.0	130.0	3504.0	12.0	70	
1	15.0	8	350.0	165.0	3693.0	11.5	70	
2	18.0	8	318.0	150.0	3436.0	11.0	70	
3	16.0	8	304.0	150.0	3433.0	12.0	70	
4	17.0	8	302.0	140.0	3449.0	10.5	70	
...	...	...	...	...	...	...	...	...
393	27.0	4	140.0	86.0	2790.0	15.6	82	
394	44.0	4	97.0	52.0	2130.0	24.6	82	
395	32.0	4	135.0	84.0	2295.0	11.6	82	
396	28.0	4	120.0	79.0	2625.0	18.6	82	
397	31.0	4	119.0	82.0	2720.0	19.4	82	

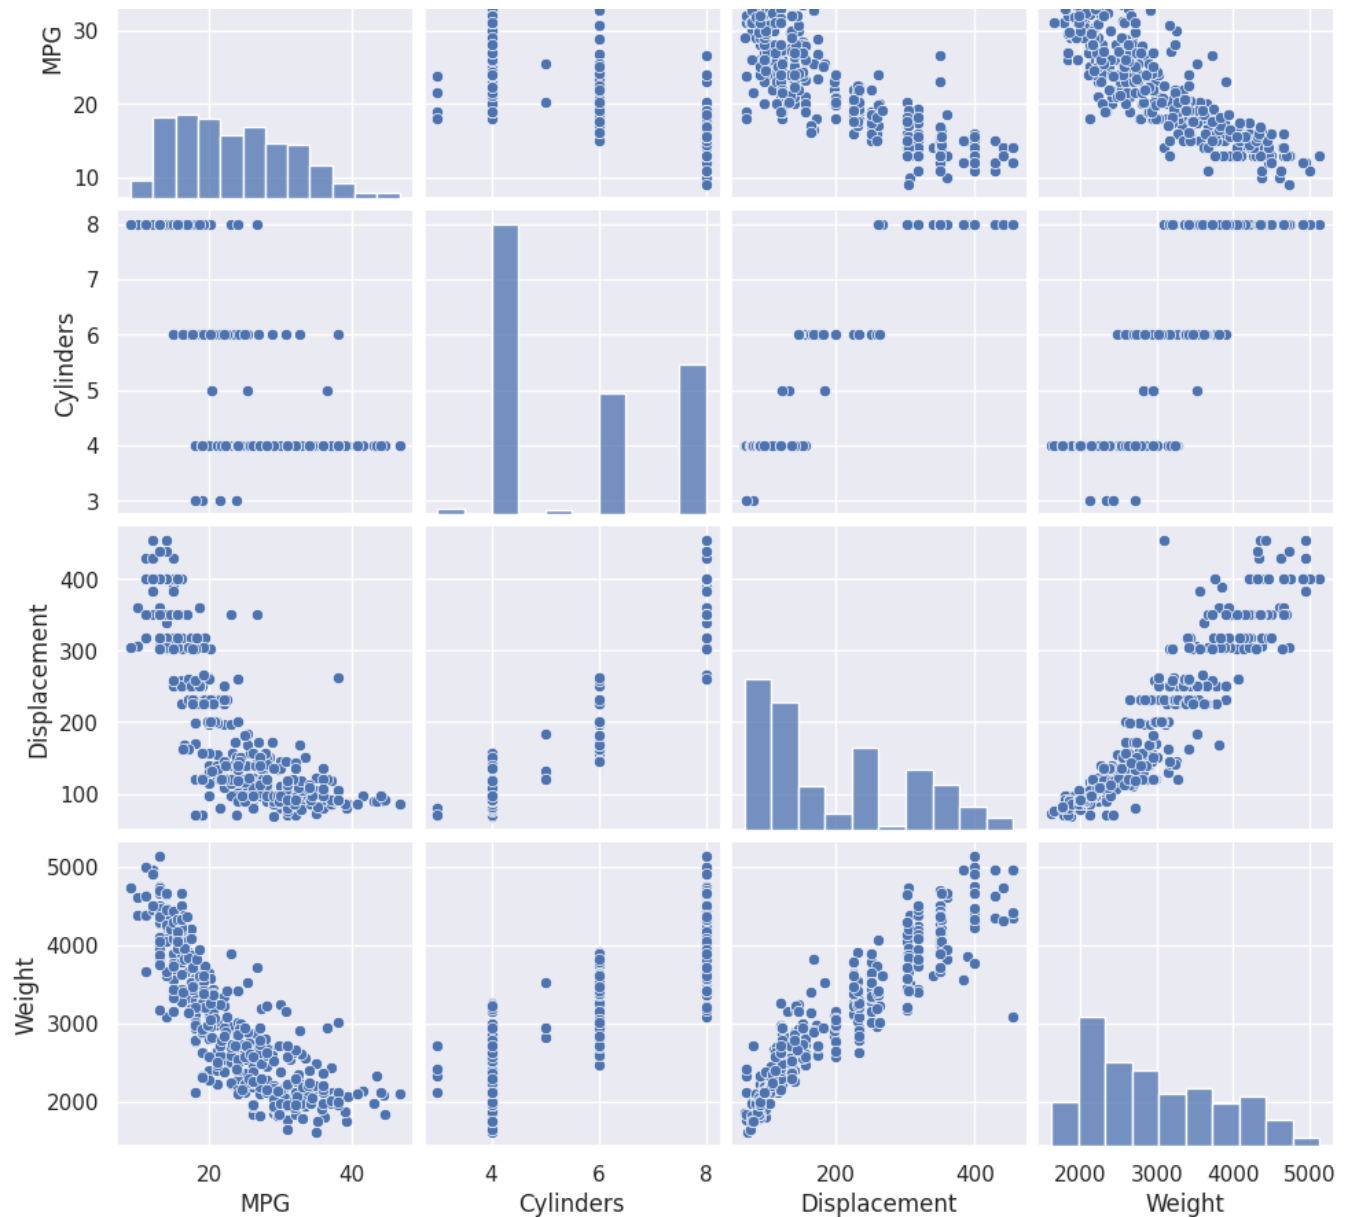
392 rows × 10 columns

```
# Uncomment to Inspect Training Data
```

```
sns.pairplot(dataset[['MPG', 'Cylinders', 'Displacement', 'Weight']])  
dataset.describe().transpose()
```

	count	mean	std	min	25%	50%	75%	
MPG	392.0	23.445918	7.805007	9.0	17.000	22.75	29.000	
Cylinders	392.0	5.471939	1.705783	3.0	4.000	4.00	8.000	
Displacement	392.0	194.411990	104.644004	68.0	105.000	151.00	275.750	4
Horsepower	392.0	104.469388	38.491160	46.0	75.000	93.50	126.000	2
Weight	392.0	2977.584184	849.402560	1613.0	2225.250	2803.50	3614.750	51
Acceleration	392.0	15.541327	2.758864	8.0	13.775	15.50	17.025	
Model Year	392.0	75.979592	3.683737	70.0	73.000	76.00	79.000	
Europe	392.0	0.173469	0.379136	0.0	0.000	0.00	0.000	
Japan	392.0	0.201531	0.401656	0.0	0.000	0.00	0.000	
USA	392.0	0.625000	0.484742	0.0	0.000	1.00	1.000	





**1.2.1.** Normalize the dataset based on the minimum and maximum value of each feature ("Displacement", "Horsepower", "Weight", "Acceleration") and label ("MPG"). Divide the original dataset into the training set (70%) and a testing set (30%).



## #1.2.1 to complete

```

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MinMaxScaler

All_Data = dataset.to_numpy()

# Normalize features. see: https://scikit-learn.org/stable/modules/generated/sk
#scaler =
#All_Data_normalized =

#Create X matrix and Y vector
Y = All_Data_normalized[:, 0:1] #Label: "MPG"
X = All_Data_normalized[:, 2:6] #Selecting only the features "Displacement, Horsepower, Weight, Acceleration"
feature_custom = np.add(X[:, 1], X[:, 3]).reshape(-1, 1) # Add feature: Horsepower + Acceleration
X = np.concatenate((X, feature_custom), axis = 1 )
# Feature vector: "Displacement", "Horsepower", "Weight", "Acceleration", "Horsepower + Acceleration"

# Split data into train and test set.
#X_train, X_test, y_train, y_test = train_test_split(...)

```

```

-----
NameError                                Traceback (most recent call last)
<ipython-input-10-7b313f305416> in <cell line: 13>()
    11
    12 #Create X matrix and Y vector
--> 13 Y = All_Data_normalized[:, 0:1] #Label: "MPG"
    14 X = All_Data_normalized[:, 2:6] #Selecting only the features
"Displacement, Horsepower, Weight, Acceleration"
    15 feature_custom = np.add(X[:, 1], X[:, 3]).reshape(-1, 1) # Add
feature: HorsePower + Acceleration

NameError: name 'All Data normalized' is not defined

```

### #Solution 1.2.1

```

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MinMaxScaler

All_Data = dataset.to_numpy()

# Normalize features. see: https://scikit-learn.org/stable/modules/generated/sk
scaler = MinMaxScaler()
All_Data_normalized = scaler.fit_transform(All_Data)

#Create X matrix and Y vector
Y = All_Data_normalized[:, 0:1] #Label: "MPG"
X = All_Data_normalized[:, 2:6] #Selecting only the features "Displacement, t
feature_custom = np.add(X[:, 1], X[:, 3]).reshape(-1, 1) # Add feature: HorseP
X = np.concatenate((X, feature_custom), axis = 1 )
# Feature vector: "Displacement", "Horsepower", "Weight", "Acceleration", "Hors

# Split data into train and test set.
X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=0.3, random

```

### 1.2.2. Determine a solution for the regression model :

$$y = a_0 + a_1x_1 + a_2x_2 + a_3x_3 + a_4x_4 + a_5x_5$$

where  $x_1, x_2, x_3, x_4$  and  $x_5$  is the displacement, horsepower, weight, acceleration and horsepower+acceleration of a car, respectively. The solution that was obtained is unique?

# 1.2.2 to complete

```
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error

#model =
#model.fit(...)
#yfit = model.predict(...)

# The mean squared error
print("Training MSE (Linear Model): %.2f" % mean_squared_error(y_train, model.p
print("Testing MSE (Linear Model): %.2f" % mean_squared_error(y_test, yfit))

# Show model
print("Model: ", model.intercept_, model.coef_)

#Singular values
print("Model Singular values:", model.singular_)
```

# Solution 1.2.2

```
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error

model = LinearRegression(fit_intercept=True)
model.fit(X_train, y_train)
yfit = model.predict(X_test)

# The mean squared error
print("Training MSE (Linear Model): %.2f" % mean_squared_error(y_train, model.p
print("Testing MSE (Linear Model): %.2f" % mean_squared_error(y_test, yfit))

# Show model
print("Model: ", model.intercept_, model.coef_)

#Singular values
print("Model Singular values:", model.singular_)
```

```
Training MSE (Linear Model): 0.01
Testing MSE (Linear Model): 0.01
Model: [0.7] [[ 9.740e-03  1.445e+12 -5.906e-01  1.445e+12 -1.445e+12]]
Model Singular values: [6.738 2.815 1.365 0.904 0.  ]
```

**1.2.3.** Implement a Ridge Regressor ( $X^T X + \lambda I$ ) and determine a solution for the previous problem. Explain whether  $X^T X + \lambda I$  is invertible.

# 1.2.3 to complete

```
from sklearn.linear_model import Ridge
from sklearn.metrics import mean_squared_error

#Ridge Model
#model_ridge =
#model_ridge.fit(...)
#yfit_ridge = model_ridge.predict(...)

# The mean squared error
print("Training MSE (Ridge Model): %.2f" % mean_squared_error(y_train, model_ri
print("Testing MSE (Ridge Model): %.2f" % mean_squared_error(y_test, yfit_ridge

# Show model
print("Model: ", model_ridge.intercept_, model_ridge.coef_)
```

# Solution 1.2.3

```
from sklearn.linear_model import Ridge
from sklearn.metrics import mean_squared_error

#Ridge Model
model_ridge = Ridge(alpha=0.1)
model_ridge.fit(X_train, y_train)
yfit_ridge = model_ridge.predict(X_test)

# The mean squared error
print("Training MSE (Ridge Model): %.2f" % mean_squared_error(y_train, model_ri
print("Testing MSE (Ridge Model): %.2f" % mean_squared_error(y_test, yfit_ridge

# Show model
print("Model: ", model_ridge.intercept_, model_ridge.coef_)
```

```
Training MSE (Ridge Model): 0.01
Testing MSE (Ridge Model): 0.01
Model: [0.704] [[-0.018 -0.14 -0.554 0.041 -0.099]]
```

**1.2.4** Repeat the previous exercise, but using LASSO regularization.

# 1.2.4 to complete ...

### # Solution 1.2.4

```

from sklearn.linear_model import Lasso
from sklearn.metrics import mean_squared_error

#Ridge Model
model_lasso = Lasso(alpha=0.0001)
model_lasso.fit(X_train, y_train)
yfit_lasso = model_lasso.predict(X_test)

# The mean squared error
print("Training MSE (LASSO Model): %.2f" % mean_squared_error(y_train, model_la
print("Testing MSE (LASSO Model): %.2f" % mean_squared_error(y_test, yfit_lassc

# Show model
print("Model: ", model_lasso.intercept_, model_lasso.coef_)

    Training MSE (LASSO Model): 0.01
    Testing MSE (LASSO Model): 0.01
    Model: [0.697] [-0.    -0.173 -0.587  0.    -0.043]

```

---

## ✓ Activity 1.3

Derive the least squares optimal solution with ridge regularization for the model

$$f_{\theta}(x) = \theta^{\top} \phi(x)$$

where  $\theta = [\theta_0, \theta_1]^{\top}$ ,  $\phi(x) = [\phi_0(x), \phi_1(x)]^{\top}$ , and  $x \in \mathbb{R}^d$ . In other words, apply the (necessary) optimality condition to

$$\min_{\theta} \left( \sum_{n=1}^N (y_n - \theta^{\top} \phi(x_n))^2 + \lambda \|\theta\|^2 \right)$$

## Solution

Let's consider a more general case where  $\theta = [\theta_0, \dots, \theta_{K-1}]$  and  $\phi(x) = [\phi_0(x), \dots, \phi_{K-1}(x)]^\top$  so that the linear regression with ridge regularization corresponds to obtain  $\hat{\theta}$  that minimizes  $J(\theta)$  defined by

$$J(\theta) = \sum_{n=1}^N (y_n - \theta^\top \phi(x_n))^2 + \lambda \|\theta\|^2$$

Defining

$$Y := \begin{bmatrix} y_1 \\ y_2 \\ \dots \\ y_N \end{bmatrix}$$

and

$$\Phi := \begin{bmatrix} \phi(x_1)^\top \\ \phi(x_2)^\top \\ \dots \\ \phi(x_n)^\top \end{bmatrix} = \begin{bmatrix} \phi_0(x_1) & \phi_1(x_1) & \phi_2(x_1) & \dots & \phi_{K-1}(x_1) \\ \phi_0(x_2) & \phi_1(x_2) & \phi_2(x_2) & \dots & \phi_{K-1}(x_2) \\ \dots & \dots & \dots & \dots & \dots \\ \phi_0(x_n) & \phi_1(x_n) & \phi_2(x_n) & \dots & \phi_{K-1}(x_n) \end{bmatrix}$$

we have

$$J(\theta) = (Y - \Phi\theta)^\top (Y - \Phi\theta) + \lambda \theta^\top \theta = Y^\top Y - 2Y^\top \Phi\theta + \theta^\top \Phi^\top \Phi\theta + \lambda \theta^\top \theta$$

and, therefore,

$$\nabla J = -2\Phi^\top Y + 2\Phi^\top \Phi\theta + 2\lambda\theta.$$

At the minimum,  $\nabla J = 0$ , resulting in  $(\Phi^\top \Phi + \lambda I)\hat{\theta} = \Phi^\top Y$ .

As  $\Phi^\top \Phi$  is at least semi-definite positive,  $\Phi^\top \Phi + \lambda I$  is non-singular, so

$$\hat{\theta} = (\Phi^\top \Phi + \lambda I)^{-1} \Phi^\top Y$$

## ✓ Part 2 - Estimation

Given  $N$  observations,  $\mathcal{X} = \{x_1, x_2, \dots, x_N\}$ , drawn from a probability distribution.

Assume that the joint pdf of these  $N$  observations is of a known parametric functional type, denoted as  $p(\mathcal{X}; \theta)$ , where the parameter  $\theta \in \mathbb{R}^K$  is unknown. The task is to estimate its value.

## ✓ Maximum Likelihood Estimation

The joint pdf,  $p(\mathcal{X}; \theta)$ , is known as the **likelihood function** of  $\theta$  with respect to the given set of observations,  $\mathcal{X}$ . According to the maximum likelihood method, the estimate is provided by

$$\hat{\theta}_{\text{ML}} = \arg \max_{\theta} p(\mathcal{X}; \theta).$$

Since the logarithmic function is monotone and increasing, one can instead search for the maximum of the **log-likelihood function**, that is,

$$\left. \frac{\partial \ln p(\mathcal{X}; \theta)}{\partial \theta} \right|_{\theta=\hat{\theta}_{\text{ML}}} = 0.$$

## ✓ Activity 2.1

The exponential distribution with parameter  $\lambda > 0$  is characterized by the probability function

$$p(x; \lambda) = \lambda e^{-\lambda x}, \quad x \geq 0.$$

Obtain the maximum likelihood estimate of  $\lambda$  based on observations  $\{x_i\}_{i=1}^N$  of this distribution.

## Solution

Assuming that observations are independent, we have

$$p(\mathcal{X}; \lambda) = \prod_{i=1}^N p(x_i; \lambda) = \prod_{i=1}^N (\lambda e^{-\lambda x_i})$$

and

$$\ln p(\mathcal{X}; \lambda) = \sum_{i=1}^N (\ln \lambda - \lambda x_i) = N \ln \lambda - \lambda \sum_{i=1}^N x_i.$$

Therefore,

$$\frac{d \ln p(\mathcal{X}; \lambda)}{d\lambda} = \frac{N}{\lambda} - \sum_{i=1}^N x_i$$

The maximum likelihood estimate,  $\lambda_{\text{ML}}$ , satisfies

$$\frac{N}{\lambda_{\text{ML}}} - \sum_{i=1}^N x_i = 0$$

or

$$\lambda_{\text{ML}} = \frac{N}{\sum_{i=1}^N x_i}.$$

## ✓ Maximum A-Posteriori Probability Estimation

The Maximum A-Posteriori Probability estimation technique, usually denoted as MAP, is based on the Bayesian theorem, but it does not go as far as the Bayesian philosophy allows to. The goal becomes that of obtaining an estimate by maximizing

$$\begin{aligned} \hat{\theta}_{\text{MAP}} &= \arg \max_{\theta} p(\theta | \mathcal{X}) = \arg \max_{\theta} \frac{p(\mathcal{X} | \theta) p(\theta)}{p(\mathcal{X})} \\ \hat{\theta}_{\text{MAP}} &= \arg \max_{\theta} p(\theta | \mathcal{X}) = \arg \max_{\theta} \frac{p(\mathcal{X} | \theta) p(\theta)}{p(\mathcal{X})}. \end{aligned}$$

Since  $p(\mathcal{X})$  is independent of  $\theta$ , this leads to

$$\begin{aligned} \hat{\theta}_{\text{MAP}} &= \arg \max_{\theta} p(\mathcal{X} | \theta) p(\theta) = \arg \max_{\theta} \{\ln p(\mathcal{X} | \theta) + \ln p(\theta)\} \\ \hat{\theta}_{\text{MAP}} &= \arg \max_{\theta} p(\mathcal{X} | \theta) p(\theta) = \arg \max_{\theta} \{\ln p(\mathcal{X} | \theta) + \ln p(\theta)\}. \end{aligned}$$



## ✓ Activitiy 2.2

Assume that  $\{x_1, x_2, \dots, x_N\}$  are i.i.d. observations from a Gaussian distribution  $\mathcal{N}(\mu, \sigma^2)$ . Obtain the MAP estimate of  $\mu$  if the prior follows the exponential distribution  $p(\mu) = \lambda e^{-\lambda\mu}$ ,  $\lambda > 0$ ,  $\mu \geq 0$ .

### Solution

The pdf of the Gaussian distribution  $p(x|\mu)$  is  $p(x|\mu) = \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{1}{2}\left(\frac{x-\mu}{\sigma}\right)^2}$  so  $p(\mathcal{X}|\mu) = \prod_{i=1}^N \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{1}{2}\left(\frac{x_i-\mu}{\sigma}\right)^2}$  and  $\ln p(\mathcal{X}|\mu) = \sum_{i=1}^N \left(-\ln\left(\frac{1}{\sqrt{2\pi}\sigma}\right) - \frac{1}{2}\left(\frac{x_i-\mu}{\sigma}\right)^2\right)$  and  $\ln p(\mu) = \ln\lambda - \lambda\mu$ .  $\mu_{\text{MAP}}$  maximizes  $M(\mu) = \ln p(\mathcal{X}|\mu) + \ln p(\mu) = -N \ln\left(\frac{1}{\sqrt{2\pi}\sigma}\right) - \frac{1}{2\sigma^2} \sum_{i=1}^N (x_i - \mu)^2 + \ln\lambda - \lambda\mu$ . The derivative is  $\frac{dM}{d\mu} = \sum_{i=1}^N \frac{x_i - \mu}{\sigma^2} - \lambda = -\lambda + \frac{1}{\sigma^2} \left(\sum_{i=1}^N x_i - N\mu\right)$ . so,  $\mu_{\text{MAP}}$  is given by  $\mu_{\text{MAP}} = \frac{\left(\sum_{i=1}^N x_i\right) - \sigma^2\lambda}{N}$ .

